



UNIVERSITÉ LIBRE DE BRUXELLES

# TPC-DS Benchmarking with Postgres Database

Data Warehouses (INFO-H413)

Erasmus Mundus Joint Master's Degree  
in  
Big Data Management and Analytics

by

Ahmad (ahmad@ulb.be)  
Rishika Gupta (rishika.gupta@ulb.be)  
Mir Wise Khan (mir.khan@ulb.be)  
Chidiebere Ogbuchi (chidiebere.ogbuchi@ulb.be)

under the guidance of

**Prof. Esteban Zimanyi**



École polytechnique de Bruxelles  
Université Libre de Bruxelles  
October 2022

# Contents

<b>Abstract</b>	<b>i</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Overview . . . . .	1
1.2 Aim and Objectives . . . . .	1
1.3 Tools Used . . . . .	1
1.4 Limitations and Justifications . . . . .	2
<b>2 Technology Fundamentals</b>	<b>3</b>
2.1 Postgres SQL . . . . .	3
2.1.1 Why PostgreSQL . . . . .	4
2.2 Other Tools . . . . .	4
2.2.1 Docker Desktop . . . . .	4
2.2.2 Visual Studio Code . . . . .	4
2.2.3 Python Interpreter . . . . .	5
<b>3 Benchmarking and Implementation</b>	<b>6</b>
3.1 Introduction to Benchmarking . . . . .	6
3.2 TPC-DS . . . . .	6
3.3 Implementing TPC-DS on PostgreSQL . . . . .	7
3.3.1 Scaling and Database Population . . . . .	8
3.3.1.1 Scaling Model . . . . .	8
3.3.1.2 Test Database Scaling . . . . .	9
3.3.1.3 DSDGEN and Database Population . . . . .	9
3.3.2 Queries Overview . . . . .	10
3.3.2.1 Query Definition . . . . .	10
3.3.2.2 Query Modifications . . . . .	10
3.3.2.3 Query Ordering . . . . .	11
3.3.3 Query Execution and Optimisation . . . . .	11
3.3.3.1 Query Execution . . . . .	11
3.3.3.2 Optimisation . . . . .	11
<b>4 Results and Discussions</b>	<b>12</b>
4.1 Benchmarking results . . . . .	12
4.1.1 Overall performance across all scale factors . . . . .	12
4.1.2 Individual performance across all scale factors . . . . .	13
4.1.3 Comparison of optimized and original queries . . . . .	15
4.2 Analysis of irregular pattern results . . . . .	15
4.2.1 Demystifying the performance variation on Query 04 . . . . .	15

4.2.2	Understanding the catalyst for queries with exponential run time . . . . .	18
4.3	Summary of Results . . . . .	20
<b>5</b>	<b>Conclusion</b>	<b>21</b>
	<b>Bibliography</b>	<b>i</b>
<b>A</b>	<b>Appendix</b>	<b>iii</b>
A.1	Python Scripts . . . . .	iii
A.2	Indexes . . . . .	xi
A.3	Optimised Queries . . . . .	xii

# List of Figures

3.1	TPC-DS Process Diagram . . . . .	7
3.2	Initial Query For Interval Syntax . . . . .	10
3.3	Updated Query For Interval Syntax . . . . .	10
4.1	Total Average Runtime (99 Queries) by Scale Factor . . . . .	12
4.2	[Query 01-20] Average Runtime (Mins) by Scale Factor (SF) . . . . .	13
4.3	[Query 21-40] Average Runtime (Mins) by Scale Factor (SF) . . . . .	13
4.4	[Query 41-60] Average Runtime (Mins) by Scale Factor (SF) . . . . .	13
4.5	[Query 61-80] Average Runtime (Mins) by Scale Factor (SF) . . . . .	13
4.6	[Query 81-99] Average Runtime (Mins) by Scale Factor (SF) . . . . .	14
4.7	Queries without exponential runtime (majority) . . . . .	14
4.8	Optimised Queries Average Runtime (Mins) by Scale Factor (SF) . . . . .	15
4.9	Original Queries Average Runtime (Mins) by Scale Factor (SF) . . . . .	15
4.10	Query 04 - Average Runtime (Mins) by Scale Factor (SF) . . . . .	15
4.11	Query 04 - Query Plan 5 SF . . . . .	16
4.12	Query 04 - Query Plan 15 SF . . . . .	17
4.13	Query 04 - Query Statistics 5 SF . . . . .	17
4.14	Query 04 - Query Statistics 15 SF . . . . .	18
4.15	Exponential run-time detailed analysis for Q14, Q30, Q95 - All SFs . . . . .	19
4.16	Query 14 - Query Statistics . . . . .	19
4.17	Query 23 - Query Statistics . . . . .	19

# List of Tables

1.1	Tools Used for TPC-DS Benchmarking with Postgres SQL . . . . .	2
3.1	Local Machine Specifications . . . . .	8
3.2	Implemented Scale Factors . . . . .	8
3.3	Tuple Count Summary . . . . .	9

# Listings

A.1	Preprocessing_DBSetup_DataLoad Script . . . . .	iii
A.2	Query_Run_Test Script . . . . .	vi
A.3	Query_Performance_Test Script . . . . .	viii
A.4	Index_Setup Script . . . . .	xi
A.5	Query_04 . . . . .	xii
A.6	Query_04 . . . . .	xii
A.7	Query_06 . . . . .	xv
A.8	Query_11 . . . . .	xvi
A.9	Query_74 . . . . .	xviii
A.10	Query_81 . . . . .	xx

# List of Abbreviations

CTE	Common Table Expressions
DB	Database
DBMS	Database Management Systems
DS	Decision Support
GB	Gigabytes
MS	Milli-seconds
OLAP	Online Analytical Processing
OLTP	Online Transaction Processing
ORDBMS	Object Relational Database Management System
RDBMS	Relational Database Management Systems
SF	Scale Factor
SQL	Structured Query Language
SUT	System Under Test
TPC	Transactional Processing Council
TPC-DS	Transactional Processing Council's Decision Support Benchmark
WSL	Windows Subsystem for Linux

## **Abstract**

TPD-DS is a profound decision support benchmark presently developed by the Transaction Processing Performance Council (TPC). It comprises elements that can be utilized to evaluate a wide spectrum of implementation methodologies mapped to a conventional business setting. This report establishes a concept that briefly outlines the business modeling systems and performance aspects adopted into this benchmark. The TPC-DS provides decision support functions of a retail product supplier, as well as data loading, query generation and data maintenance. The database contains several snowflake schemas with common tabling dimensions with a huge bag of queries. In general, the benchmarks model very essential aspects of a typical decision support system which entails transformation of transactional data into business intelligence as well as synchronization and maintenance processes of data structures.



# Chapter 1

## Introduction

### 1.1 Overview

This project entails selecting a suitable Database Management System (DBMS) tool on which the TPC-DS benchmark will be implemented such as SQL Server Analysis Services, SparkSql, PostgreSQL, MariaDB, etc.

The benchmark is performed with several scale factors which relatively influence the data warehouse size. A reference scale factors is estimated, including other factors at different dimensions in order to evaluate its performance.

The project is performed in a groups of 4 persons and delivers a self-explanatory report of the major significant parts of the implementation. This report employs PostgreSQL as the preferred choice of DBMS on which the TPC-DS benchmark is performed.

### 1.2 Aim and Objectives

- The main aim of this report is to implement and evaluate the TPC-DS benchmark on a DBMS tool in order to learn how to efficiently perform a benchmark.
- Illustrating the objective further, with this project we will understand and implement the TPC-DS benchmark on a DBMS tool, learn how to perform a benchmark that helps us choose the best available database tool for our business.
- Another objective is to evaluate the benchmarking performance and analyse the results.

### 1.3 Tools Used

The tools installed and utilized to perform the benchmark operation are summarized in the table 1.1.

Tool	Version	Description
TPC-DS standard benchmark tool	3.2.0	The official tools set offered by TPC-DS for data generation, query generation and an answer set to compare results.
PostgreSQL	14.0	Open-source PostgreSQL relational database to store and query data.
JupyterLab	3.3.2	IDE for Python and iPython notebook.
Docker Desktop	20.10.17	Docker was used to run Ubuntu on top of Windows to run the TPC-DS tool in a Linux environment.
Python	3.10.5	Python was used as the main programming language for running our scripts.
MS PowerBI	2.110	Power BI was utilized to create visualizations of benchmarking results for further analysis.
GitHub	2.38.1	GitHub Desktop was used to share code files as well as images conveniently with the team members.

Table 1.1: Tools Used for TPC-DS Benchmarking with Postgres SQL

## 1.4 Limitations and Justifications

Since the entire project was implemented on a local machine, it was associated with a certain limit on the resources (tools that could be used), thereby preventing scaling higher benchmarks. It is definitely possible for us to procure cloud-based services like Google Cloud, and Amazon Azure and implement Postgres in their environments. With the provision of more cores as well as storage, we could have certainly benchmarked until 100 GB's least. Despite the complexity associated with local resources, we tried to benchmark 20 GB as well as 25 GB but haven't included them in the report due to abstract results.

# Chapter 2

## Technology Fundamentals

### 2.1 Postgres SQL

PostgreSQL is a free enterprise open-source object relational database management system (ORDBMS) akin to a relational database, but that it is object-oriented such that it offers classes and objects models including inheritance in query-language and database schemas [Bartolini et al., 2017]. Initially developed at the University of California, Berkeley by the Database Research Team of the computer science department, is now adapted and developed by a vast horde of contributory developers. It provides a huge diversity of support languages ranging from C, Python, PHP, C++, Perl and Java amongst others that permits a variety selection of constructs that can proffer solutions to problems [The PostgreSQL Global Development Group, 2022]. In benchmarks, PostgreSQL is fast and provides similar excellent performance as when compared to other proprietary and open source databases [Obe & Hsu, 2017]. Also, it shoulders a huge part of the SQL standard and offers advanced present-day features such as but not limited to:

- Complex queries
- Transactional integrity Triggers
- Multiversion concurrency control
- Foreign keys
- Updatable views [Matthew & Stones, 2005]

Furthermore, PostgreSQL allows user extension in several ways such as adding and connecting new:

- operators
- data types
- index methods
- procedural languages
- aggregate functions
- functions

As a result of the open license, PostgreSQL can be utilized, distributed & modified by any individuals without charge for any reason [The PostgreSQL Global Development Group, 2022].

### 2.1.1 Why PostgreSQL

PostgreSQL has numerous benefits including:

- Outstanding SQL standards compliance.
- Client-server architectural structure.
- High degree of synchronous interface and design where users don't interfere with each other.
- High extent of configuration and extensions for several kinds of applications
- Outstanding scalability and performance with high-level tuning and optimization features.
- Excellent support for different types of data formats including relational, post-relational (arrays, nested relations via record types) documents (JSON, CSV and XML), and dictionary keys/values.

In addition, the PostgreSQL system is a robust and high-quality tool with rich documentation, maintainability, interoperability and high availability. It requires low maintenance as well as provides excellent performance, security and compatibility for major operating systems on both enterprise and embedded usage [Bartolini et al., 2017]. In this project, PostgreSQL shall be used as a database management tool for implementing the TPC-DS benchmark.

## 2.2 Other Tools

### 2.2.1 Docker Desktop

This application allows for the transformation and optimization of workflows by allowing users to connect to a collection of pre-built developer tools and systems from the Docker Extension Marketplace. It allows for the creation and sharing of customized tools with other team members in its dev environment.

Also, Docker provides a fast way to build solutions and projects in containers as well as offers flexible control, secure access and management of container images [Install Docker Desktop, 2022]. For this coursework, Docker was used as a replacement to WSL (Windows Subsystem for Linux) to run the latest version of TPC-DS tool on Linux (Ubuntu) for the purpose of generating executable SQL queries for PostgreSQL from query templates.

### 2.2.2 Visual Studio Code

Visual Studio Code is a compact but extremely powerful source code editor that runs on computer desktops and is accessible on macOS, Windows and Linux operating systems. It has a built-in interface standard for Typescript, Node.js and JavaScript as well as a offers a wide array of extensions for other programming languages (Python, C++, C, Java, etc.). In action, visual studio code has an impressive UX and allows the customization of workflows [Visual Studio Code, 2022]. This project was useful for building and verifying the entire solution on dsdgen.sln.

### 2.2.3 Python Interpreter

Python is a general-purpose programming language that allows quick working and integration of systems effectively. This high-level language is dynamically input and supports procedural, functional and object-oriented programmed. It can be compiled using an interactive development emulator [[Python, 2022](#)]. For this benchmark project, Jupyter notebook was used to create and compile python scripts. Python allowed us to cleanse and transform the initial load data generated and push them into the database. Also, it was used to wrangle the generated query templates for effective accessibility on PostgreSQL.

# Chapter 3

## Benchmarking and Implementation

### 3.1 Introduction to Benchmarking

Benchmarking involves comparing performance indicators and processes to industry best practices usually in relation to time, quality and cost metrics. It is generally used to estimate similarities and contrast between a specific performance metric. In databases, benchmarking may be difficult especially if it follows different relational and object model approach. Despite this fact, organizations and individuals still experience the challenge of selecting a suitable DBMS platform for implementing models, as most databases offer many similar features on many fronts. However, performance is a great differentiator when choosing between available databases for decision support. Leveraging benchmarks can be used in recommending a suitable selection of a given technology [[Tortosa, 2020](#)].

In other words, benchmarking a database is the process of performing well-defined tests on that particular database for the purpose of evaluating its performance [[Kabangu, 2009](#)]. The performance evaluation can help an organization decide if the particular choice of the database can meet the business needs of the organization in the long run.

### 3.2 TPC-DS

TPC Benchmark™DS (TPC-DS) is basically a decision support benchmark model that fashions various relevant areas of a simple decision support structure, entailing queries as well as data maintenance. The TPC-DS benchmark offers a comprehensive decision-making system that represents a typical appraisal of the System Under Test's (SUT) performance model. Generally, it depicts a typical decision support platform that:

- Evaluates huge amounts of data;
- Provides solutions to business challenges in reality;
- Performs queries that meet the need of several operational complexities such as ad-hoc, data mining and reporting requirements;

- Delineate a large CPU usage and IO loading;
- Typifies periodic database maintenance activities especially with OLTP database synchronization;
- Executes on ORDBMS and RDBMS based systems.

In addition, a benchmark result assesses various aspects including the query response time output in an isolated user level, query throughput in multiple user levels and data maintenance evaluation for a designated hardware, data processing and operating system setting under a monitored and controlled decision support workload [Transaction Processing Performance Council (TPC), 2021].

### 3.3 Implementing TPC-DS on PostgreSQL

Figure 3.1 shows a business process model depicting a brief rundown of the implementation the TPC-DS benchmark on PostgreSQL.

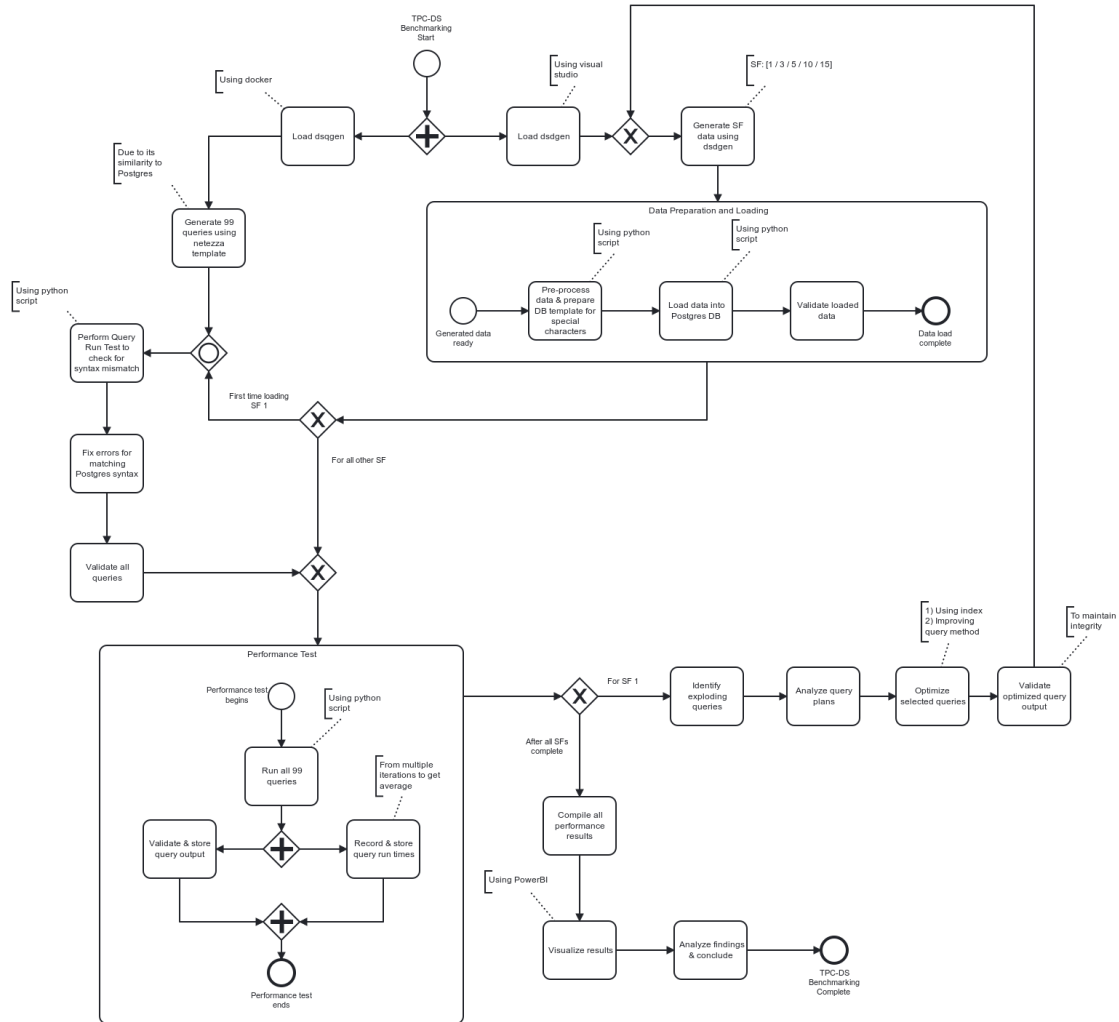


Figure 3.1: TPC-DS Process Diagram

Some of the major processes undertaken are summarized as follows:

- Building the solution of the dsdgen and dsqgen from the TPC-DS tool kit using visual studio and docker desktop respectively to obtain input data for the initial loading of the warehouse via the dsdgen utility.
- Generating scale factors using the dsdgen for the load data as well as generating 99 queries via the dsqgen using the Netezza SQL template.
- Cleaning, wrangling and transforming the data using suitable python commands to prepare and load the data into the database.
- Matching the queries with PostgreSQL by performing a query run test to check for mismatches and error.
- Using python scripts, perform tests on the queries and compile all results.
- Discover exploding queries in order to analyze and optimize them.
- Re-run all scale factors and evaluate final results.
- Visualize all outputs and provide detailed reports.

The benchmark was implemented on a local machine with specifications illustrated in table 3.1:

CPU (AMD Ryzen 7 6800HS)	RAM (DDR5 SODIMM)	GPU (NVIDIA GeForce RTX3060)
<ul style="list-style-type: none"> <li>• 8 cores, 16 threads</li> <li>• Base clocking speed at 3.2GHz and can over-clock up to 4.7GHz</li> <li>• 16MB L3 Cache</li> </ul>	<ul style="list-style-type: none"> <li>• 16 GB memory</li> <li>• 4800MHz speed</li> </ul>	<ul style="list-style-type: none"> <li>• Dedicated graphics</li> <li>• 6GB VRAM</li> </ul>

Table 3.1: Local Machine Specifications

### 3.3.1 Scaling and Database Population

#### 3.3.1.1 Scaling Model

TPC-DS Benchmark identifies a set of distinct points used for scaling, hence called the scale factors that depend on the dsdgen file (from the TPC-DS toolkit), which is influenced by the software and hardware on which it is run [Transaction Processing Performance Council (TPC), 2021].

Scale Factors set for this project are listed in the table 3.2, wherein gigabyte (GB) is equivalent to  $2^{10}$  bytes.

SF	1	3	5	10	15
Scale Factor	1 GB	3 GB	5 GB	10 GB	15 GB

Table 3.2: Implemented Scale Factors



As each scale factor has a corresponding SF which has no units and is almost equal to the bytes stored in the database. For this project the various scale factors and SFs are presented in the table 3.2

### 3.3.1.2 Test Database Scaling

For each scale factor, the total number of tuples stored in the data warehouse differs. This variation is depicted as in the table 3.3.

DB Table: Tuple Count Summary (for each scale factor)					
Table Name	SF 1	SF 3	SF 5	SF 10	SF 15
call_center	6	10	14	24	6
catalog_page	11718	11718	11718	12000	11718
catalog_returns	144067	432000	720174	1439749	2160757
catalog_sales	1441548	4319367	7199490	14401261	21602679
customer	100000	188000	277000	500000	183000
customer_address	50000	94000	138000	250000	91000
customer_demographics	1920800	1920800	1920800	1920800	1920800
date_dim	73049	73049	73049	73049	73049
household_demographics	7200	7200	7200	7200	7200
income_band	20	20	20	20	20
inventory	11745000	28188000	49329000	133110000	14356305
item	18000	36000	54000	102000	22000
promotion	300	344	388	500	327
reason	35	37	39	45	35
ship_mode	20	20	20	20	20
store	12	32	52	102	28
store_returns	287514	862834	1437911	2875432	4315222
store_sales	2880404	8639377	14400052	28800991	43197400
time_dim	86400	86400	86400	86400	86400
warehouse	5	6	7	10	5
web_page	60	90	122	200	162
web_returns	71763	215477	359991	719217	1079028
web_sales	719384	2160165	3599503	7197566	10795812
web_site	30	32	34	42	30

Table 3.3: Tuple Count Summary

### 3.3.1.3 DSDGEN and Database Population

The data is loaded into the database using the following command:

```
.\dsdgen.exe /scale 10 /dir .\tmp /suffix .csv /parallel 4 /child 1
/delimiter "^" /terminate n &
```

```
.\dsdgen.exe /scale 10 /dir .\tmp /suffix .csv /parallel 4 /child 2
/delimiter "^" /terminate n &
```

```
.\dsdgen.exe /scale 10 /dir .\tmp /suffix .csv /parallel 4 /child 3
/delimiter "^" /terminate n &
```

```
.\dsdgen.exe /scale 10 /dir .\tmp /suffix .csv /parallel 4 /child 4
/delimiter "^" /terminate n
```

- Scale factor varied as per the table 3.2
- Data was dumped as CSV file (Default file output extension is .dat)
- For fast and parallel data generation, the process is run parallelly on 4 threads
- Each child process command is concatenated using and
- End of line (EOL) is indicated by “|” and if “|” appears consecutively twice, it indicates a NULL value in the column. This is handled using:
  - “^” as the delimiter -> /delimiter “^”
  - End of line (EOL) character to be NULL -> /terminate n
- To support special international characters throughout the entire workflow, ISO/IEC 8859-1 is used. For Postgres - the Win1252 template supports the requisite character sets [[The PostgreSQL Global Development Group, 2022](#)]

### 3.3.2 Queries Overview

#### 3.3.2.1 Query Definition

Queries in TPC-DS basically answer various business questions, pertaining to the data warehouse for example - “What are the total sales through each channel in the year 2009?” TPC-DS toolkit includes several templates for query generation. In this project, we generated queries as per the Netezza template due to its close resemblance with PostgreSQL.

#### 3.3.2.2 Query Modifications

To ensure complete compatibility of queries with PostgreSQL, a few minor modifications were done.

- Intervals: Postgres requires the specific keyword - interval to identify the interval of days, Query Syntax were modified from 3.2 to 3.3.

```
and d_date between (cast ('2002-05-18' as date) - 30 days)
and (cast ('2002-05-18' as date) + 30 days)
```

Figure 3.2: Initial Query For Interval Syntax

```
and d_date between (cast ('2002-05-18' as date) - interval '30 day')
and (cast ('2002-05-18' as date) + interval '30 day')
```

Figure 3.3: Updated Query For Interval Syntax

- Aliases - A few of the sub-queries and columns were given an alias “x”.
- Joins - Instead of generally selecting from Table A, Table B and specifying the JOIN condition in WHERE Clause, explicit joins (inner/left/right) were used.

### 3.3.2.3 Query Ordering

Queries are ordered in sequences by the dsqgen file. Due to the difference in the query sequences, -QUALIFY was used as the sample command below to generate them in the same order as the answer templates.

```
./dsqgen -DIRECTORY ../query_templates -INPUT
../query_templates/templates.lst \ -VERBOSE Y -QUALIFY Y
-DIALECT netezza
```

*Note: The above command was executed on a Docker (Linux) environment.*

## 3.3.3 Query Execution and Optimisation

### 3.3.3.1 Query Execution

- After query generation by TPC-DS templates, queries were run to identify the syntax errors. Modifications to the queries were done as in 3.3.2.2 for 23 erroneous queries in this project.
- Queries were first run on SF-1 and the average run-time for 5 iterations was considered as the Query run-time.
- The previous step was essential to identify the set of potential queries that might explode (the run-time of that particular query is very high in comparison to the others) as data increases (higher SFs) because of correlated sub-queries, sub-queries, etc.
- Optimisation (as in 3.5.2) was done for a few queries and run against all SFs.

### 3.3.3.2 Optimisation

- Optimized queries are detailed in Appendix A.
- The ways used in this project to optimize are:
  - Query plans were studied thoroughly to analyze possible points of optimisation.
  - Correlated Sub-queries were rewritten using Common Table Expressions (CTE).
  - Indexes were added for fast retrieval of table tuples.
  - Tables in joins were re-ordered in the increasing order of their counts ensuring efficient joins.
  - Distinct keyword was eliminated by rewriting queries in the form of Common Table Expressions (CTE).

# Chapter 4

## Results and Discussions

As mentioned in the previous section, the benchmark was performed on a local machine, with a total of five different scale factors (1, 3, 5, 10 and 15). The results will be focusing on a few different criteria, namely, the evolution of the run times as the scale factor of the database increases, comparison between optimized and original query performances, analysis of queries with exponential run time growth, and the SQL methods that typically hampers the performance of the query execution time.

### 4.1 Benchmarking results

#### 4.1.1 Overall performance across all scale factors

Inspecting the total average run time for all the 99 queries which were run sequentially for multiple iterations, it is evident that the performance is nowhere close to being linear. In fact, as the scale factor increases further, an exponential pattern starts to emerge, as seen in Figure 4.1.




Figure 4.1: Total Average Runtime (99 Queries) by Scale Factor

At a scale factor (SF) of 1 (equivalent to a database size of 1GB), all of the 99 queries were managed to complete within a span of 5 minutes (on average). This quickly rises to a total of 58 minutes for SF 3, which then follows along with a

smaller increase in run time for SF 5 and 10 respectively, but rises again steeply when progressing to SF 15.

Although this is the case, looking only at the overall run time for all 99 queries together, gives a biased assumption regarding the performance of many of the individual queries. Thus, the performance results are further broken down into individual queries in the next section.

#### 4.1.2 Individual performance across all scale factors

Due to a large number of queries, they are split into five charts containing 20 queries each, as seen in Figures 4.2 to 4.6. Firstly, it can be observed that the majority of the queries are not running exponentially as the scale factor increases and it is actually due to a few specific queries such as query 14, 30 and 95 that causes the total overall run time to massively increase as seen previously in Figure 4.1. Furthermore, an unusual pattern is evident with query 4, where it initially rises in run time for SF 3 and 5, but reduces as the scale factor increases more towards SF 10 and 15. Both of these unique observations are further discussed in Section 4.2.




Figure 4.2: [Query 01-20] Average Runtime (Mins) by Scale Factor (SF)




Figure 4.3: [Query 21-40] Average Runtime (Mins) by Scale Factor (SF)




Figure 4.4: [Query 41-60] Average Runtime (Mins) by Scale Factor (SF)




Figure 4.5: [Query 61-80] Average Runtime (Mins) by Scale Factor (SF)




Figure 4.6: [Query 81-99] Average Runtime (Mins) by Scale Factor (SF)

Given that there were a few queries with exponential run times, the scale for the charts above was stretched due to these outliers, resulting in less visibility on the patterns for some of the queries with much less run time. Therefore, an additional chart is generated below (Figure 4.7), which excludes these outliers and shows the overall non-exponential trend for the majority of the queries (95%), and mostly managing to complete within 60 seconds at a scale factor of 15.




Figure 4.7: Queries without exponential runtime (majority)

### 4.1.3 Comparison of optimized and original queries




Figure 4.8: Optimised Queries Average Runtime (Mins) by Scale Factor (SF)




Figure 4.9: Original Queries Average Runtime (Mins) by Scale Factor (SF)

As discussed in the benchmarking implementation (Section 3), a total of six queries (1, 4, 6, 11, 74 and 81) were optimized by improving the structure and SQL method-/approach used (while maintaining the same output as original). From Figure 4.8 (optimized queries), it is evident that none of the six optimized queries continued to explode exponentially as the scale factors progressed up to SF 15. Although, as mentioned previously, query 4 experienced a unique pattern which is discussed in Section 4.2. In contrast, if we observe Figure 4.9 (original queries), it can be seen that there are a few queries (14, 30 and 95) that kept increasing exponentially in run time as the scale factor increased (also discussed in Section 4.2). Nonetheless, most of the remaining queries did not explode possibly due to the help of indices as well.

## 4.2 Analysis of irregular pattern results

### 4.2.1 Demystifying the performance variation on Query 04

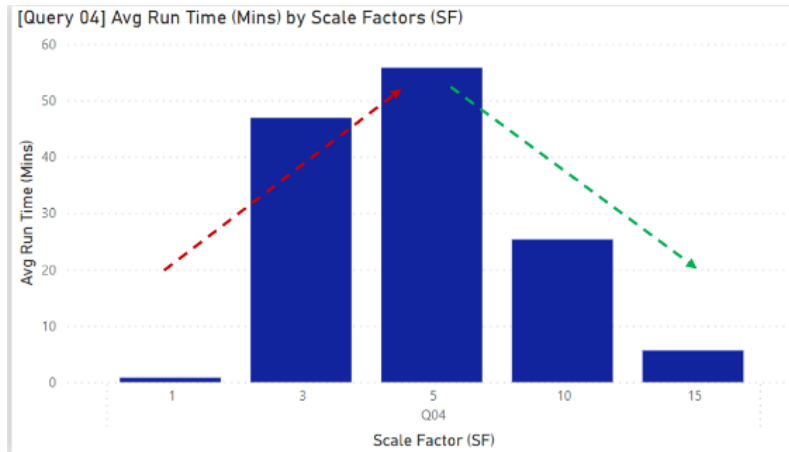


Figure 4.10: Query 04 - Average Runtime (Mins) by Scale Factor (SF)

As mentioned above, query 04 was one of the queries that were optimized to further improve its performance as the database is scaled. This is due to the query making

use of all three large fact tables (catalog\_sales, store\_sales, web\_sales), which can result in very expensive steps for the query to run. As seen from Figure 4.10, the query worked very well for SF 1, but as the scale factor increased to 3 and 5, the run time increased exponentially, but then interestingly reduced when scaling to SF 10 and 15.

In order to understand what changed between SF 5 and SF 15, a deeper dive was done into the query planner and statistics in terms of how the query was executed in steps for both of these scale factors.

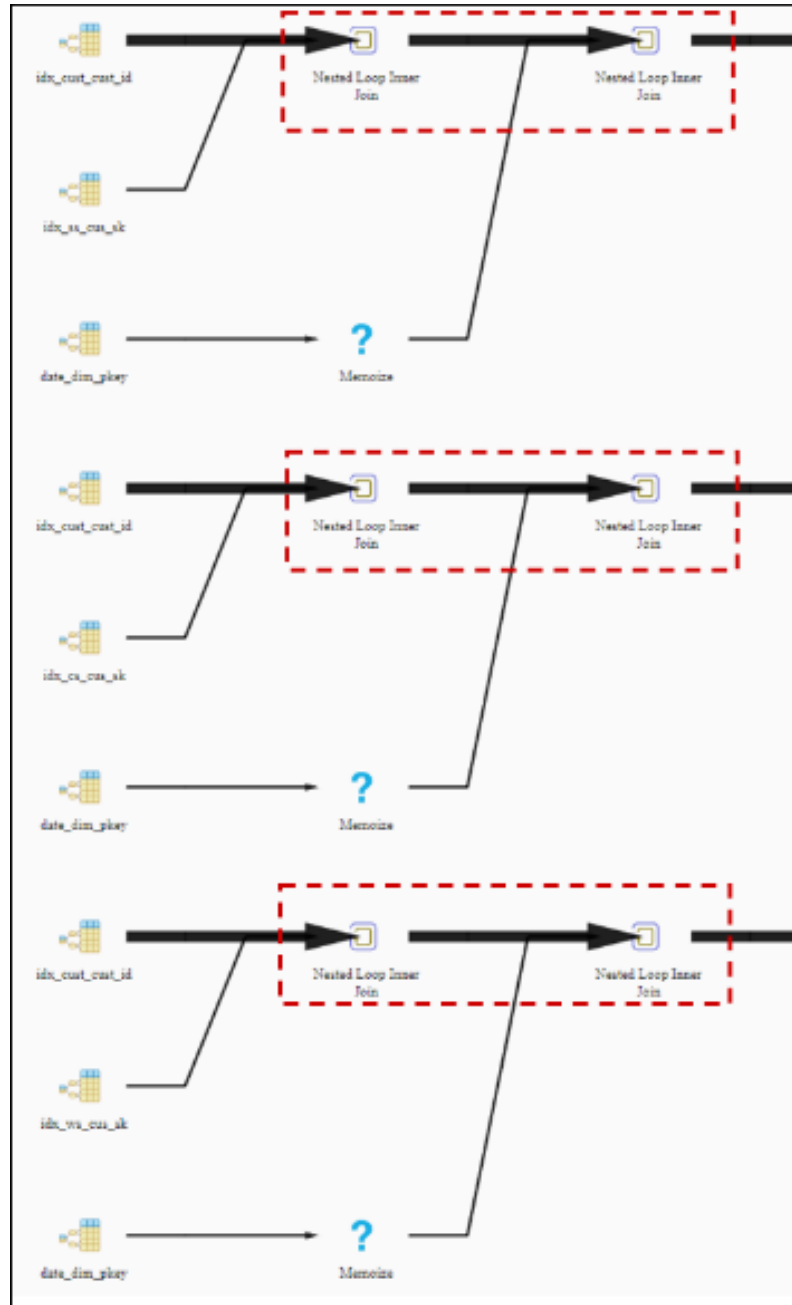


Figure 4.11: Query 04 - Query Plan 5 SF



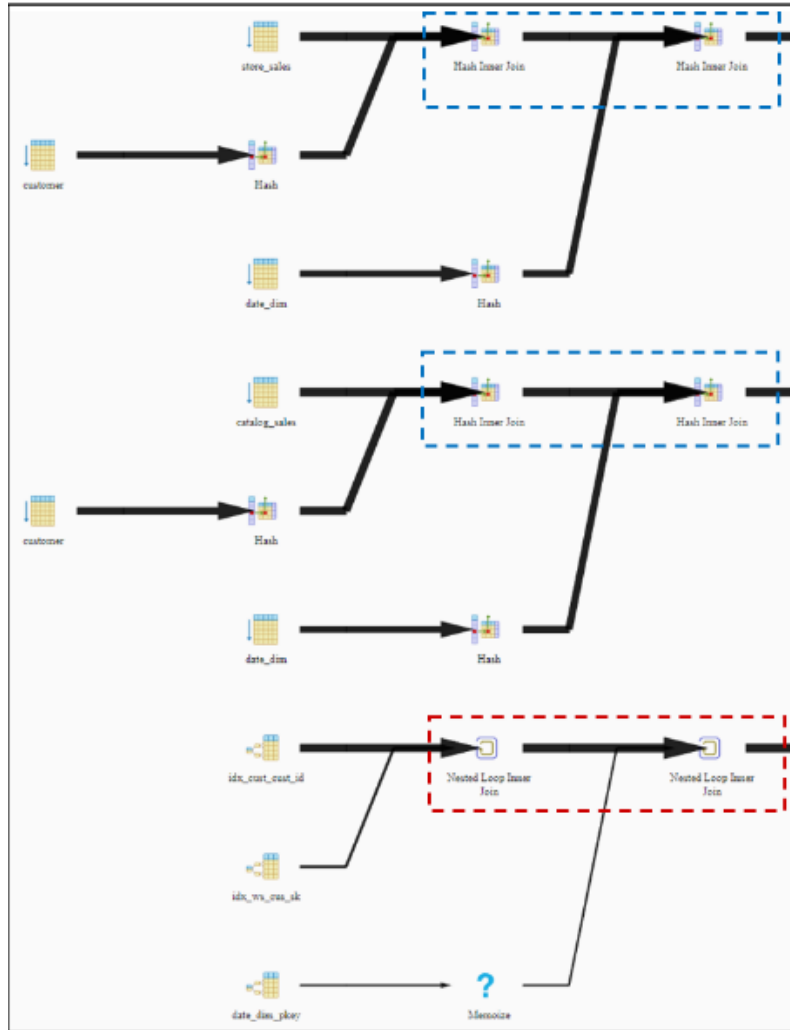


Figure 4.12: Query 04 - Query Plan 15 SF

Statistics per Node Type				Statistics per Relation			
Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query
Aggregate	6	7355.196 ms	0.21%	Node type	Count	Sum of times	% of relation
CTE Scan	6	111459.894 ms	3.12%	catalog_sales	1	0.05 ms	0.01%
Gather Merge	3	47002.046 ms	1.32%	Index Scan	1	0.05 ms	100%
Incremental Sort	4	9329.735 ms	0.27%	customer	3	571.742 ms	0.02%
Index Scan	9	571.891 ms	0.02%	Index Scan	3	571.742 ms	100%
Limit	1	3471878.941 ms	96.94%	date_dim	3	0.003 ms	0.01%
Materialize	2	15.91 ms	0.01%	Index Scan	3	0.003 ms	100%
Memoize	3	0 ms	0%	store_sales	1	0.072 ms	0.01%
Merge Inner Join	4	397.699 ms	0.02%	Index Scan	1	0.072 ms	100%
Nested Loop Inner Join	7	3514598.933 ms	98.13%	web_sales	1	0.024 ms	0.01%
Sort	5	551.103 ms	0.02%	Index Scan	1	0.024 ms	100%

Figure 4.13: Query 04 - Query Statistics 5 SF

Statistics per Node Type				Statistics per Relation			
Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query
Aggregate	6	17101.377 ms	4.59%	Node type	Count	Sum of times	% of relation
CTE Scan	6	101151.402 ms	27.13%	catalog_sales	1	3653.605 ms	0.98%
Gather Merge	3	34629.793 ms	9.29%	Seq Scan	1	3653.605 ms	100%
Hash	4	234.046 ms	0.07%	customer	3	265.619 ms	0.08%
Hash Inner Join	4	21792.748 ms	5.85%	Index Scan	1	193.142 ms	72.72%
Incremental Sort	2	4976.232 ms	1.34%	Seq Scan	2	72.477 ms	27.29%
Index Scan	3	193.253 ms	0.06%	date_dim	3	7.408 ms	0.01%
Limit	1	131542.595 ms	35.28%	Index Scan	1	0.001 ms	0.02%
Materialize	3	0.78 ms	0.01%	Seq Scan	2	7.407 ms	99.99%
Memoize	1	0 ms	0%	store_sales	1	1412 ms	0.38%
Merge Inner Join	3	4.698 ms	0.01%	Seq Scan	1	1412 ms	100%
Nested Loop Inner Join	4	293517.387 ms	78.71%	web_sales	1	0.11 ms	0.01%
Seq Scan	6	5145.489 ms	1.38%	Index Scan	1	0.11 ms	100%
Sort	6	135467.997 ms	36.33%				

Figure 4.14: Query 04 - Query Statistics 15 SF

Firstly, inspecting the query plan for both scale factors, it's evident that different initial steps were taken. As seen from Figure 4.11 for SF 5, the query planner decided to make use of all indexes available, but then proceeded to the next step with multiple nested loop inner joins. On the other hand, in Figure 4.13 for SF 15, the query planner felt that the customer table was small enough in comparison to two out of three of the large fact tables (catalog\_sales store\_sales) and decided to perform a sequential scan for them instead of using an index. This resulted in less usage of nested loop inner joins overall for the next steps.

Looking further into the statistics for the query run, it can be observed that the additional nested loop inner joins resulted in a much more expensive execution overall. Thus, for SF 5, the multiple nested loop inner joins cost a total of 58 minutes (3514598.933 ms), while for SF 15, it cost only about 5 minutes (293517.387 ms) as there were much fewer of them (specifically avoiding the large fact tables). Detailed statistics are available in Figures 4.13 and 4.14.

This observation further confirms that the index can play the hero and the devil in different situations, given that the query plan's choices could result in more expensive decisions for the steps that follow that.

## 4.2.2 Understanding the catalyst for queries with exponential run time

Three queries demonstrated exponential run times, namely query 14, 30 and 95 (Figure 4.15), and this creates the opportunity to further understand what were the factors that influenced this the most.

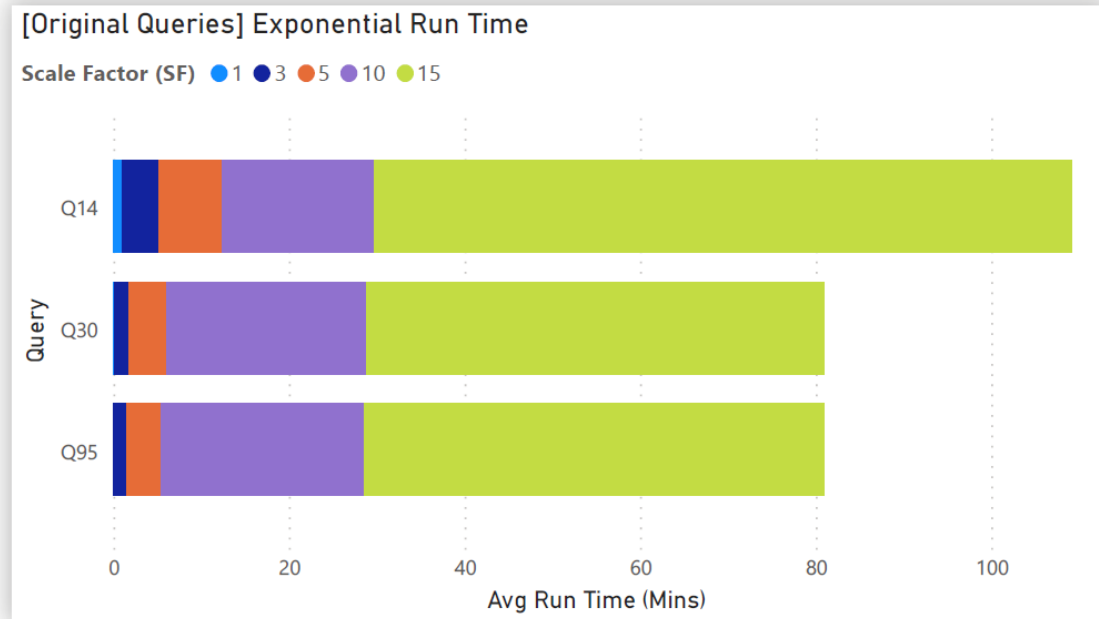


Figure 4.15: Exponential run-time detailed analysis for Q14, Q30, Q95 - All SFs

To begin with, we have query 14, which is a unique query that actually has two outputs. In order to make a fair comparison, query 23 was used as it also has two outputs and uses all the three large fact tables as well. Similar to the previous analysis, a deeper dive was done into the query planner and statistics, which resulted in an identical discovery where excessive usage of nested loop inner joins on large tables for query 14 (Figure 4.16) resulted in a much longer run time (exponential), as compared to query 23 (Figure 4.17) which had more of a linear run time as the scale factors progressed.

Statistics per Node Type				Statistics per Relation				Statistics per Node Type				Statistics per Relation			
Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query	Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query
Aggregate	9	3076.298 ms	0.41%	Node type	Count	Sum of times	% of relation	Aggregate	6	2465.649 ms	0.25%	Node type	Count	Sum of times	% of query
Append	4	2484.09 ms	0.32%	catalog_sales	3	1905.409 ms	0.26%	Append	3	2384.906 ms	0.22%	catalog_sales	2	1881.136 ms	0.18%
Bitmap Heap Scan	1	21.22 ms	0.01%	Index Scan	1	8.822 ms	0.47%	Bitmap Heap Scan	2	37.535 ms	0.01%	Seq Scan	2	1881.136 ms	100%
Bitmap Index Scan	1	0.001 ms	0.01%	Seq Scan	2	1997.987 ms	99.54%	Bitmap Index Scan	2	0.860 ms	0.01%	data_dim	10	46.523 ms	0.01%
CTE Scan	6	27363.652 ms	3.58%	data_dim	9	1.793 ms	0.01%	CTE Scan	4	26418.017 ms	2.47%	Index Scan	10	46.523 ms	100%
Gather	4	6169.239 ms	0.81%	Index Scan	9	1.793 ms	100%	Gather	4	5923.643 ms	0.55%	item	6	20.248 ms	0.01%
Hash	13	7.764 ms	0.01%	item	7	24.547 ms	0.81%	Hash	10	4.722 ms	0.01%	Index Scan	2	0.968 ms	0.34%
Hash Inner Join	13	24667.091 ms	3.17%	Index Scan	3	0.061 ms	0.33%	Hash Inner Join	10	13378.268 ms	1.42%	Seq Scan	4	20.18 ms	99.67%
Hashed Intersect	2	5170.589 ms	0.68%	Seq Scan	4	24.466 ms	99.68%	Hashed Intersect	2	4926.412 ms	0.46%	store_sales	4	1577.421 ms	0.15%
Index Scan	14	16.466 ms	0.01%	store_sales	3	1892.449 ms	0.21%	Index Scan	12	49.591 ms	0.01%	Bitmap Heap Scan	2	37.525 ms	2.38%
Limit	1	1059950.316 ms	97.54%	Bitmap Heap Scan	1	21.22 ms	1.84%	Limit	1	1059950.316 ms	97.54%	Seq Scan	2	1529.896 ms	97.63%
Nested Loop Inner Join	6	772636.581 ms	75.07%	Seq Scan	2	1571.329 ms	99.67%	Nested Loop Inner Join	6	1349841.947 ms	129.81%	web_sales	2	2202.029 ms	0.21%
Result	1	0.537 ms	0.01%	web_sales	3	2862.498 ms	0.27%	Result	1	25.27 ms	0.00%	Seq Scan	2	2202.029 ms	100%
Seq Scan	10	9490.01 ms	0.73%	Index Scan	1	5.77 ms	0.29%	Seq Scan	10	5643.226 ms	0.52%				
Sort	4	1422.78 ms	0.19%	Seq Scan	2	1999.728 ms	99.72%	Sort	5	77203.501 ms	7.12%				
Subquery Scan	4	3224.136 ms	0.43%					Subquery Scan	4	3104.948 ms	0.29%				

Figure 4.16: Query 14 - Query Statistics

Statistics per Node Type				Statistics per Relation				Statistics per Node Type				Statistics per Relation			
Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query	Node type	Count	Time spent	% of query	Relation name	Scan count	Total time	% of query
Aggregate	12	19354.644 ms	14.91%	Node type	Count	Sum of times	% of relation	Aggregate	13	14984.768 ms	15.12%	Node type	Count	Sum of times	% of relation
Append	1	0.816 ms	0.01%	catalog_sales	1	2.637 ms	0.01%	Append	1	0.04 ms	0.01%	catalog_sales	1	2.472 ms	0.01%
Bitmap Heap Scan	2	3.715 ms	0.01%	Bitmap Heap Scan	1	2.637 ms	100%	Bitmap Heap Scan	2	3.526 ms	0.01%	Bitmap Heap Scan	1	2.472 ms	100%
Bitmap Index Scan	2	0.002 ms	0.01%	customer	2	4.705 ms	0.01%	Bitmap Index Scan	2	0.002 ms	0.01%	customer	4	7.516 ms	0.01%
CTE Scan	4	84671.148 ms	82.18%	Index Only Scan	2	4.705 ms	100%	CTE Scan	4	81472.66 ms	82.21%	Index Only Scan	2	7.477 ms	99.49%
Gather Merge	3	2772.679 ms	2.7%	data_dim	4	1.292 ms	0.01%	Gather Merge	3	2804.833 ms	2.83%	Index Scan	2	0.039 ms	0.02%
Hash	9	36.637 ms	0.04%	Index Scan	4	1.292 ms	100%	Hash	9	23.889 ms	0.03%	data_dim	4	0.511 ms	0.01%
Hash Inner Join	9	15175.104 ms	14.73%	item	1	5.831 ms	0.01%	Hash Inner Join	9	14076.363 ms	14.21%	Index Scan	4	0.511 ms	100%
Index Only Scan	2	4.705 ms	0.01%	Seq Scan	1	9.831 ms	100%	Index Only Scan	2	7.477 ms	0.01%	item	1	6.48 ms	0.01%
Index Scan	4	1.292 ms	0.01%	store_sales	3	2056.076 ms	2.56%	Index Scan	6	0.55 ms	0.01%	Seq Scan	1	6.48 ms	100%
Limit	1	16862.237 ms	16.21%	Seq Scan	3	2056.076 ms	100%	Limit	1	18329.332 ms	18.26%	store_sales	3	2523.958 ms	2.55%
Nested Loop Inner Join	2	77991.501 ms	77.47%	web_sales	1	1.878 ms	0.01%	Nested Loop Inner Join	4	17360.938 ms	17.42%	Seq Scan	3	2523.958 ms	100%
Seq Scan	2	2671.607 ms	2.59%	Bitmap Heap Scan	1	1.878 ms	100%	Seq Scan	4	2530.858 ms	2.56%	web_sales	1	1.054 ms	0.01%
Sort	3	48935.991 ms	47.11%					Sort	6	46997.968 ms	47.33%	Bitmap Heap Scan	1	1.054 ms	100%

Figure 4.17: Query 23 - Query Statistics

Query 30 and 95 also demonstrated similar performance issues as seen above, and therefore will not be elaborated further.

## 4.3 Summary of Results

Running a benchmark test on a large variety of queries, on multiple scale factors provided great insights into how Postgres deals with both increasing volume of data, indices and the alteration of node steps chosen depending on the size of the table data. Although majority of the queries demonstrated non-exponential run-time performance (more than 90%), there were a few that stood out. An important influential factor was identified from these queries, which is that nested loop inner joins are one of the most taxing operations that can be done during query execution, and Postgres struggles with that once the data volume increases.

# Chapter 5

## Conclusion

In conclusion, the TPC-DS framework has provided a well-rounded model, including pragmatic components to perform a fair and transparent benchmark on Postgres DB. Some queries were optimized to further improve their performance given they were struggling in the initial run test and a few index were also added to full utilize the capabilities of the database. This enabled the performance test to have even more variety, allowing additional analysis in terms of comparison between original and optimized queries, the identification of exploding queries with exponential run time performance and understanding unique behaviours between different scale factors. A key component that has been identified, was the heavy computational cost of nested loop inner joins which caused a few of the queries to struggle and have large execution run times as the scales increased. Therefore, it is crucial for analysts to design their OLAP queries in such a way that they can avoid correlated sub-queries which results in the usage of nested loop inner joins. Furthermore, it is important to always consider how the query planner is going to form the execution steps, in order to fully utilized the capabilities of its resources and avoid settling for higher computational steps as much as possible.

# Bibliography

[Bartolini et al., 2017] Bartolini, G., Ciolli, G., & Riggs, S. (2017). PostgreSQL Administration Cookbook - Third Edition. Packt Publishing, Limited.

[Install Docker Desktop, 2022] install Docker Desktop. (2022, January 22). Docker. Retrieved October 27, 2022, from <https://www.docker.com/products/docker-desktop/>

[Kabangu, 2009] Kabangu, S. (2009). Benchmarking Databases.

[Matthew & Stones, 2005] Matthew, N., & Stones, R. (2005). Beginning Databases with PostgreSQL: From Novice to Professional. Apress.

[Obe & Hsu, 2017] Obe, R. O., & Hsu, L. S. (2017). PostgreSQL: Up and Running : a Practical Guide to the Advanced Open Source Database. In (pp. 11 - 14). O'Reilly Media, Incorporated.

[The PostgreSQL Global Development Group, 2022] The PostgreSQL Global Development Group. (2022). PostgreSQL 14.5 Documentation. PostgreSQL. Retrieved October 26, 2022, from <https://www.postgresql.org/docs/14>

[Python, 2022] Python. (2022, February). Welcome to Python.org. Retrieved October 27, 2022, from <https://www.python.org/>

[Tortosa, 2020] Tortosa, Á. H. (2020). Performance Benchmark Postgresql / Mongoddb. Ongres.com. Retrieved October 28, 2022, from [https://info.enterprisedb.com/rs/069-ALB-339/images/PostgreSQL\\_MongoDB\\_Benchmark-WhitepaperFinal.pdf](https://info.enterprisedb.com/rs/069-ALB-339/images/PostgreSQL_MongoDB_Benchmark-WhitepaperFinal.pdf)

[Transaction Processing Performance Council (TPC), 2021] TPC-DS. (2021, January). tpc.org. Retrieved October 28, 2022, from <https://www.tpc.org/tpcds/>

[Visual Studio Code, 2022] Visual Studio Code. (2022, January). Visual Studio Code - Code Editing. Redefined. Retrieved October 27, 2022, from

<https://code.visualstudio.com/>

[Zimanyi, n.d.] Zimanyi, E. (n.d.). INFO-H-419: Data Warehouses. INFO-H-419: Data Warehouses [Université Libre de Bruxelles - Service CoDE - Laboratoire WIT]. Retrieved October 27, 2022, from <https://cs.ulb.ac.be/public/teaching/infoh419>

# Appendix A

## Appendix

### A.1 Python Scripts

```
1 # TPCDS: Preprocessing, DB Setup and Data Load Script
2
3 # importing Libraries
4 import sys, os, re
5 import psycopg2
6 import numpy as np
7 import pandas as pd
8 from psycopg2 import Error
9 from psycopg2.extensions import ISOLATION_LEVEL_AUTOCOMMIT
10
11 # set up connection variables
12 db_host = "localhost"
13 db_port = "5432"
14 db_user = "postgres"
15 db_pass = "password"
16 db_name = "postgres"
17
18 # function to connect with postgres
19 def connect_postgres(db_host, db_port, db_user, db_pass, db_name):
20     try:
21         # Connect to an existing database
22         connection = psycopg2.connect(host = db_host,
23                                     port = db_port,
24                                     user = db_user,
25                                     password = db_pass,
26                                     database = db_name)
27
28         # Set auto-commit
29         connection.set_isolation_level(ISOLATION_LEVEL_AUTOCOMMIT);
30         # Create a cursor to perform database operations
31         cur = connection.cursor()
32         # Print PostgreSQL details
33         print("PostgreSQL server information")
34         print(connection.get_dsn_parameters(), "\n")
35         # Executing a SQL query
36         cur.execute("SELECT version();")
37         # Fetch result
38         record = cur.fetchone()
39         print("You are connected to - ", record, "\n")
40     except (Exception, Error) as error:
```



```

41         print("Error while connecting to PostgreSQL", error)
42     else:
43         return cur
44
45 # connect to postgres
46 cur = connect_postgres(db_host, db_port, db_user, db_pass, db_name)
47
48 # drop tpcds db
49 db_name = "tpcds"
50
51 cur.execute(
52     f"DROP DATABASE IF EXISTS {db_name} WITH (FORCE);"
53 )
54 print("SQL Status Output:\n", cur.statusmessage)
55
56
57 # change win1252 encoding temp db to normal before drop
58 try:
59     cur.execute(
60         "ALTER DATABASE win1252_temp is_template false;"
61     )
62 except Exception as e:
63     print(e)
64 else:
65     print("SQL Status Output:\n", cur.statusmessage)
66
67 # drop win1252 encoding temp db (after set to normal db)
68 cur.execute(
69     "DROP DATABASE IF EXISTS win1252_temp WITH (FORCE);"
70 )
71 print("SQL Status Output:\n", cur.statusmessage)
72
73 # create win1252 encoding temp db
74 cur.execute(
75     """
76
77     CREATE DATABASE win1252_temp
78     WITH
79     OWNER = postgres
80     TEMPLATE = template0
81     ENCODING = 'WIN1252'
82     CONNECTION LIMIT = -1
83     IS_TEMPLATE = True;
84
85     """
86 )
87 print("SQL Status Output:\n", cur.statusmessage)
88
89 # create tpcds db
90 cur.execute(
91     f"""
92
93     CREATE DATABASE {db_name}
94     WITH
95     OWNER = postgres
96     TEMPLATE = win1252_temp
97     ENCODING = 'WIN1252'
98     CONNECTION LIMIT = -1

```

```

99         IS_TEMPLATE = False;
100
101     """
102 )
103 print("SQL Status Output:\n", cur.statusmessage)
104
105 # connect to tpcds db
106 cur = connect_postgres(db_host, db_port, db_user, db_pass, db_name)
107
108 # create tables for db
109 cur.execute(open("tools/tpcds.sql", "r").read())
110 print("SQL Status Output:\n", cur.statusmessage)
111 cur.execute(open("tools/tpcds_source.sql", "r").read())
112 print("SQL Status Output:\n", cur.statusmessage)
113
114 # get dir path
115 path = os.getcwd() + '\\tools\\tmp\\'
116 files = os.listdir(path)
117 print(path)
118
119 # function to get full absolute path of csv files containing data
120 def get_absolute_path(d):
121     return [os.path.join(d, f) for f in os.listdir(d)]
122
123 # get full absolute path of csv files containing data
124 files_abs_path = [p.replace('\\', '/') for p in get_absolute_path(
    path)]
125 print("Total files:", len(files_abs_path))
126 print("First few files...")
127 files_abs_path[:5]
128
129 # exclude extra delimiter for dbgen_version file
130 file_count = 0
131 for iteration in range(0, 1):
132     for file in files_abs_path:
133         file_open = open(file, 'r')
134         all_text = file_open.read().replace(" ", "")
135         file_open.close()
136
137         if (all_text[-13] == '^' and 'dbgen_version' in file):
138             file_open_read = open(file, 'r', encoding = 'latin-1')
139             string_list = file_open_read.readlines()
140             file_open_read.close()
141
142             for i in range(len(string_list)):
143                 last_delimiter_index = string_list[i].rfind("^")
144                 string_list[i] = string_list[i][:
                    last_delimiter_index] + "" + string_list[i][last_delimiter_index
                        + 1:]
145
146                 file_open_write = open(file, 'w', encoding = 'latin-1')
147                 new_file_contents = ''.join(string_list)
148                 file_open_write.write(new_file_contents)
149                 file_open_write.close()
150
151                 file_count += 1
152     else:
153         pass

```

```

154     print(f'\nIteration {iteration + 1} done!')
155     print(f'{file_count} file(s) updated for extra column exclusion
156     .')
157     file_count = 0
158
159 # generate sql commands for loading data from csv to postgres db
160 # considers that csv files were generated in parallel stream
161 sql_commands_file = open('data_load_script.sql', 'w')
162
163 for file in files:
164     underscore_index = [underscore_ind.start() for underscore_ind
165     in re.finditer('_', file)]
166     file_name = file[:underscore_index[-2]]
167     file_path = path+file
168     sql_command = "COPY public."+file_name+" FROM '"+file_path+"
169     delimiter '^' CSV;\n"
170     sql_commands_file.write(sql_command)
171
172 sql_commands_file.close()
173
174 # load csv data into db
175 cur.execute(open("data_load_script.sql", "r").read())
176 print("SQL Status Output:\n", cur.statusmessage)
177
178 # add constraints to db
179 cur.execute(open("tools/tpcds_ri.sql", "r").read())
180 print("SQL Status Output:\n", cur.statusmessage)
181
182 # close connection to db
183 cur.close()
184
185 # End of script.

```

Listing A.1: Preprocessing\_DBSetup\_DataLoad Script

```

1 # TPCDS: Query Run Test Script
2
3 # importing libraries
4 import sys, os
5 import psycopg2
6 from psycopg2 import Error
7 from psycopg2.extensions import ISOLATION_LEVEL_AUTOCOMMIT
8 from IPython.display import clear_output
9 from datetime import datetime
10
11 # set up connection variables
12 db_host = "localhost"
13 db_port = "5432"
14 db_user = "postgres"
15 db_pass = "password"
16 db_name = "tpcds"
17
18 # function to connect with postgres
19 def connect_postgres(db_host, db_port, db_user, db_pass, db_name):
20     try:
21         # Connect to an existing database
22         connection = psycopg2.connect(host = db_host,
23                                     port = db_port,
24                                     user = db_user,

```

```

25         password = db_pass,
26         database = db_name)
27
28     # Set auto-commit
29     connection.set_isolation_level(ISOLATION_LEVEL_AUTOCOMMIT);
30     # Create a cursor to perform database operations
31     cur = connection.cursor()
32     # Print PostgreSQL details
33     print("PostgreSQL server information")
34     print(connection.get_dsn_parameters(), "\n")
35     # Executing a SQL query
36     cur.execute("SELECT version();")
37     # Fetch result
38     record = cur.fetchone()
39     print("You are connected to - ", record, "\n")
40
41     except (Exception, Error) as error:
42         print("Error while connecting to PostgreSQL", error)
43     else:
44         return cur
45
46 # connect to postgres
47 cur = connect_postgres(db_host, db_port, db_user, db_pass, db_name)
48
49 # get dir path
50 path = os.getcwd() + '\\all_queries\\initial_queries'
51 files = os.listdir(path)
52 print(path)
53
54 # function to get full absolute path files in directory
55 def get_absolute_path(d):
56     return [os.path.join(d, f) for f in os.listdir(d)]
57
58 # get full absolute path files in directory
59 files_abs_path = [p.replace('\\', '/') for p in get_absolute_path(
60     path)]
61 print("Total files:", len(files_abs_path))
62 print("First few files...")
63 files_abs_path[:5]
64
65 # printing start datetime
66 now = datetime.now()
67 current_time = now.strftime("%H:%M:%S")
68 print("Run Test Start =", current_time)
69
70 # perform run test on each query
71 # save results in text file
72 script_num = 1
73 script_errors = 0
74 for sql_script in files_abs_path:
75     textfile = open("query_run_test_result.txt", "a")
76     textfile2 = open("query_run_test_query_errors.txt", "a")
77     clear_output(wait = True)
78     try:
79         cur.execute(
80             open(sql_script, "r").read()
81         )
82     except Exception as e:
83         script_errors += 1

```

```

82         outcome = f"Error, Message: {e}"
83         print(sql_script)
84         print(outcome)
85         textfile.write(sql_script + "\n")
86         textfile.write(outcome + "\n\n")
87         # for tracking errors alone
88         textfile2.write(sql_script + "\n")
89         textfile2.write(outcome + "\n\n")
90     else:
91         outcome = f"Success, Message: {cur.statusmessage}"
92         print(sql_script)
93         print(outcome)
94         textfile.write(sql_script + "\n")
95         textfile.write(outcome + "\n\n")
96
97     script_num += 1
98     textfile.close()
99     textfile2.close()
100
101
102 # printing end datetime
103 now = datetime.now()
104 current_time = now.strftime("%H:%M:%S")
105 print("Run Test End =", current_time)
106
107 # close connection to db
108 cur.close()
109
110 # check total amount of query errors
111 print(f"We have a total of {script_errors} queries with error")
112
113 # End of script.

```

Listing A.2: Query\_Run\_Test Script

```

1 # TPCDS: Query Performance Test Script
2
3 # importing libraries
4 import sys, os
5 import psycopg2
6 import numpy as np
7 import pandas as pd
8 from psycopg2 import Error
9 from psycopg2.extensions import ISOLATION_LEVEL_AUTOCOMMIT
10 from datetime import datetime
11 from IPython.display import clear_output
12
13 # scale factor being tested
14 sf = 'sf_1'
15
16 # set up connection variables
17 db_host = "localhost"
18 db_port = "5432"
19 db_user = "postgres"
20 db_pass = "password"
21 db_name = "tpcds"
22
23 # function to connect with postgres
24 def connect_postgres(db_host, db_port, db_user, db_pass, db_name):

```

```

25     try:
26         # Connect to an existing database
27         connection = psycopg2.connect(host = db_host,
28                                       port = db_port,
29                                       user = db_user,
30                                       password = db_pass,
31                                       database = db_name)
32
33         # Set auto-commit
34         connection.set_isolation_level(ISOLATION_LEVEL_AUTOCOMMIT);
35         # Create a cursor to perform database operations
36         cur = connection.cursor()
37         # Print PostgreSQL details
38         print("PostgreSQL server information")
39         print(connection.get_dsn_parameters(), "\n")
40         # Executing a SQL query
41         cur.execute("SELECT version();")
42         # Fetch result
43         record = cur.fetchone()
44         print("You are connected to - ", record, "\n")
45
46     except (Exception, Error) as error:
47         print("Error while connecting to PostgreSQL", error)
48     else:
49         return cur
50
51 # connect to postgres
52 cur = connect_postgres(db_host, db_port, db_user, db_pass, db_name)
53
54 # get dir path
55 path = os.getcwd() + '\\all_queries\\optimized_queries_final'
56 files = os.listdir(path)
57 print(path)
58
59 # function to get full absolute path files in directory
60 def get_absolute_path(d):
61     return [os.path.join(d, f) for f in os.listdir(d)]
62
63 # get full absolute path files in directory
64 files_abs_path = [p.replace('\\', '/') for p in get_absolute_path(
65     path)]
66 print("Total files:", len(files_abs_path))
67 print("First few files...")
68 files_abs_path[:5]
69
70 # setup dataframe for recording query execution run times
71 query_name_list = []
72 for i in range(len(files)):
73     query_name_list.append("Q" + files[i][-6:-4])
74 query_name_dict = {'query':query_name_list}
75 exec_details_df = pd.DataFrame(query_name_dict)
76
77 # get the date-time before all 99 queries have run (with iterations
78 if chosen)
79
80 run_start_default = datetime.now()
81 # dd/mm/YY H:M:S
82 run_start = run_start_default.strftime("%d/%m/%Y %H:%M:%S")
83 print("Overall Run Start:", run_start)

```

```

81
82 # run all 99 queries in sequence, and multiple iterations if chosen
83 # save query result table output
84 # save query execution run time (for all iterations)
85 q_errors = 0
86 exec_details = []
87 # choose number of iterations to run
88 n_iterations = 3
89
90 for i in range(1, n_iterations + 1):
91     clear_output(wait = True)
92     print(f'Iteration {i}\n')
93     q_index = 0
94     exec_details = []
95     iteration_start = datetime.now()
96     for sql_script in files_abs_path:
97
98         exec_start = datetime.now()
99         try:
100             cur.execute(
101                 open(sql_script, "r").read()
102             )
103         except Exception as e:
104             q_errors += 1
105             outcome = "Error"
106         else:
107             outcome = "Success"
108
109         exec_end = datetime.now()
110         exec_run_time = "{:.2f}".format((exec_end - exec_start).
total_seconds())
111         query_num = query_name_list[q_index]
112         print(f'{query_num}: Success, Execution Time: {
exec_run_time}s')
113         exec_details.append(exec_run_time)
114
115         # load table output to csv file (on first iteration only)
116         if i == 1:
117             df = pd.DataFrame(cur.fetchall(), columns = [desc[0]
for desc in cur.description])
118             df.to_csv(f'performance_test/{sf}/{query_num}.csv',
index = False)
119         else:
120             pass
121
122         q_index += 1
123
124         iteration_end = datetime.now()
125         iteration_run_time = "{:.2f}".format(((iteration_end -
iteration_start).total_seconds()) / 3600)
126         print(f'\n{sf.upper()}, Iteration {i}, Total run time for the
99 queries: {iteration_run_time}hr')
127
128         # append iteration execution details to dataframe
129         exec_details_df[f'exec_time_iter_{i}'] = np.array(exec_details)
130
131 # check total amount of query errors
132 print(f"We have a total of {q_errors} queries with error")

```

```

133
134 # get the date-time after all 99 queries have run (with iterations
    if chosen)
135 run_end_default = datetime.now()
136 # dd/mm/YY H:M:S
137 run_end = run_end_default.strftime("%d/%m/%Y %H:%M:%S")
138 print(f"Overall Run End (with {n_iterations} iterations):", run_end
    )
139
140
141 # get the total run time (in hours) for all 99 queries to complete
    (with iterations if chosen)
142 total_run_time = "{:.2f}".format(((run_end_default -
    run_start_default).total_seconds()) / 3600)
143 print(f'Total run time for the 99 queries (with {n_iterations}
    iterations): {total_run_time}hr')
144
145 # full details on query execution times (including iterations &
    average)
146 # load execution details to csv
147 exec_details_df['avg_exec_time'] = np.round(exec_details_df.iloc[:,
    1:].apply(pd.to_numeric).mean(axis = 1), 2)
148 exec_details_df.to_csv(f'performance_test/{sf}/exec_time_details_{
    sf}.csv', index = False)
149 exec_details_df
150
151 # close connection to db
152 cur.close()
153
154 # End of script.

```

Listing A.3: Query\_Performance\_Test Script

## A.2 Indexes

```

1 create index if not exists idx_cs_cus_sk
2 on public.catalog_sales
3 using hash (cs_bill_customer_sk);
4
5 create index if not exists idx_cs_sold_date_sk
6 on public.catalog_sales
7 using hash (cs_sold_date_sk);
8
9 create index if not exists idx_cust_cust_id
10 on public.customer
11 using btree (c_customer_id);
12
13 create index if not exists idx_date_dyear
14 on public.date_dim
15 using btree (d_year);
16
17 create index if not exists idx_ss_cus_sk
18 on public.store_sales
19 using hash (ss_customer_sk);
20
21 create index if not exists idx_ss_sold_date_sk

```



```

22 on public.store_sales
23 using hash (ss_sold_date_sk);
24
25 create index if not exists idx_ws_cus_sk
26 on public.web_sales
27 using hash (ws_bill_customer_sk);
28
29 create index if not exists idx_ws_sold_date_sk
30 on public.web_sales
31 using hash (ws_sold_date_sk);

```

Listing A.4: Index\_Setup Script

## A.3 Optimised Queries

```

1  -- Query 01
2
3  with customer_total_return as (
4      select sr_customer_sk as ctr_customer_sk,
5             sr_store_sk as ctr_store_sk,
6             sum(sr_return_amt_inc_tax) as ctr_total_return
7      from store_returns,
8           date_dim
9      where sr_returned_date_sk = d_date_sk
10           and d_year = 1999
11      group by sr_customer_sk,
12              sr_store_sk
13  )
14
15  , average_cust_returns as (
16      select
17          ctr_store_sk as store_sk,
18          avg(ctr_total_return) * 1.2 as ctr_avg_return
19      from customer_total_return
20      group by ctr_store_sk
21  )
22
23  select
24      c_customer_id
25  from customer_total_return
26  inner join average_cust_returns
27      on ctr_store_sk = store_sk
28  inner join store
29      on ctr_store_sk = s_store_sk
30  inner join customer
31      on ctr_customer_sk = c_customer_sk
32  where s_state = 'TN'
33  and ctr_total_return > ctr_avg_return
34  order by c_customer_id
35  limit 100;

```

Listing A.5: Query\_04

```

1  -- Query 04
2
3  with store_year_total as (
4      select c_customer_id customer_id

```

```

5      ,c_first_name customer_first_name
6      ,c_last_name customer_last_name
7      ,c_preferred_cust_flag customer_preferred_cust_flag
8      ,c_birth_country customer_birth_country
9      ,c_login customer_login
10     ,c_email_address customer_email_address
11     ,d_year dyear
12     ,sum(((ss_ext_list_price-ss_ext_wholesale_cost-
13     ss_ext_discount_amt)+ss_ext_sales_price)/2) year_total
14 from customer
15     ,store_sales
16     ,date_dim
17 where c_customer_sk = ss_customer_sk
18     and ss_sold_date_sk = d_date_sk
19 group by c_customer_id
20         ,c_first_name
21         ,c_last_name
22         ,c_preferred_cust_flag
23         ,c_birth_country
24         ,c_login
25         ,c_email_address
26         ,d_year
27 )
28 , catalog_year_total as (
29 select c_customer_id customer_id
30     ,c_first_name customer_first_name
31     ,c_last_name customer_last_name
32     ,c_preferred_cust_flag customer_preferred_cust_flag
33     ,c_birth_country customer_birth_country
34     ,c_login customer_login
35     ,c_email_address customer_email_address
36     ,d_year dyear
37     ,sum((((cs_ext_list_price-cs_ext_wholesale_cost-
38     cs_ext_discount_amt)+cs_ext_sales_price)/2) ) year_total
39 from customer
40     ,catalog_sales
41     ,date_dim
42 where c_customer_sk = cs_bill_customer_sk
43     and cs_sold_date_sk = d_date_sk
44 group by c_customer_id
45         ,c_first_name
46         ,c_last_name
47         ,c_preferred_cust_flag
48         ,c_birth_country
49         ,c_login
50         ,c_email_address
51         ,d_year
52 )
53 , web_year_total as (
54 select c_customer_id customer_id
55     ,c_first_name customer_first_name
56     ,c_last_name customer_last_name
57     ,c_preferred_cust_flag customer_preferred_cust_flag
58     ,c_birth_country customer_birth_country
59     ,c_login customer_login
60     ,c_email_address customer_email_address

```

```

61         ,d_year dyear
62         ,sum((((ws_ext_list_price-ws_ext_wholesale_cost-
        ws_ext_discount_amt)+ws_ext_sales_price)/2) ) year_total
63 from customer
64     ,web_sales
65     ,date_dim
66 where c_customer_sk = ws_bill_customer_sk
67        and ws_sold_date_sk = d_date_sk
68 group by c_customer_id
69          ,c_first_name
70          ,c_last_name
71          ,c_preferred_cust_flag
72          ,c_birth_country
73          ,c_login
74          ,c_email_address
75          ,d_year
76 )
77
78 , t_s_firstyear as (
79     select customer_id
80            ,customer_first_name
81            ,customer_last_name
82            ,customer_email_address
83            ,year_total
84     from store_year_total
85     where dyear = 2001
86     and year_total > 0
87 )
88
89 , t_s_secyear as (
90     select customer_id
91            ,customer_first_name
92            ,customer_last_name
93            ,customer_email_address
94            ,year_total
95     from store_year_total
96     where dyear = 2001+1
97 )
98
99 , t_c_firstyear as (
100    select customer_id
101           ,customer_first_name
102           ,customer_last_name
103           ,customer_email_address
104           ,year_total
105    from catalog_year_total
106    where dyear = 2001
107    and year_total > 0
108 )
109
110 , t_c_secyear as (
111    select customer_id
112           ,customer_first_name
113           ,customer_last_name
114           ,customer_email_address
115           ,year_total
116    from catalog_year_total
117    where dyear = 2001+1

```

```

118 )
119
120 , t_w_firstyear as (
121     select customer_id
122            ,customer_first_name
123            ,customer_last_name
124            ,customer_email_address
125            ,year_total
126     from web_year_total
127     where dyear = 2001
128     and year_total > 0
129 )
130
131 , t_w_secyear as (
132     select customer_id
133            ,customer_first_name
134            ,customer_last_name
135            ,customer_email_address
136            ,year_total
137     from web_year_total
138     where dyear = 2001+1
139 )
140
141 select
142     t_s_secyear.customer_id
143     ,t_s_secyear.customer_first_name
144     ,t_s_secyear.customer_last_name
145     ,t_s_secyear.customer_email_address
146 from t_s_firstyear
147 inner join t_s_secyear
148     on t_s_firstyear.customer_id = t_s_secyear.customer_id
149 inner join t_w_firstyear
150     on t_s_firstyear.customer_id = t_w_firstyear.customer_id
151 inner join t_w_secyear
152     on t_w_firstyear.customer_id = t_w_secyear.customer_id
153 inner join t_c_firstyear
154     on t_s_firstyear.customer_id = t_c_firstyear.customer_id
155 inner join t_c_secyear
156     on t_s_firstyear.customer_id = t_c_secyear.customer_id
157 where
158     (t_c_secyear.year_total / nullif(t_c_firstyear.year_total, 0))
159     >
160     (t_s_secyear.year_total / nullif(t_s_firstyear.year_total, 0))
161     and
162     (t_c_secyear.year_total / nullif(t_c_firstyear.year_total, 0))
163     >
164     (t_w_secyear.year_total / nullif(t_w_firstyear.year_total, 0))
165 order by t_s_secyear.customer_id
166            ,t_s_secyear.customer_first_name
167            ,t_s_secyear.customer_last_name
168            ,t_s_secyear.customer_email_address
169 limit 100;

```

Listing A.6: Query\_04

```

1 -- Query 06
2
3 with average_item_price as (
4     select

```

```

5     i_category as category,
6     avg(i_current_price) * 1.2 as avg_item_price
7 from item
8 group by i_category
9 )
10
11 select  a.ca_state state, count(*) cnt
12 from customer_address a
13 inner join customer c
14     on a.ca_address_sk = c.c_current_addr_sk
15 inner join store_sales s
16     on c.c_customer_sk = s.ss_customer_sk
17 inner join date_dim d
18     on s.ss_sold_date_sk = d.d_date_sk
19 inner join item i
20     on s.ss_item_sk = i.i_item_sk
21 inner join average_item_price aip
22     on i.i_category = aip.category
23 where
24     d.d_year = 1998
25     and d.d_moy = 3
26     and i.i_current_price > aip.avg_item_price
27 group by a.ca_state
28 having count(*) >= 10
29 order by cnt, a.ca_state
30 limit 100;

```

Listing A.7: Query\_06

```

1  -- Query 11
2
3  with store_year_total as (
4      select  c_customer_id customer_id
5              ,c_first_name customer_first_name
6              ,c_last_name customer_last_name
7              ,c_preferred_cust_flag customer_preferred_cust_flag
8              ,c_birth_country customer_birth_country
9              ,c_login customer_login
10             ,c_email_address customer_email_address
11             ,d_year dyear
12             ,sum(ss_ext_list_price-ss_ext_discount_amt) year_total
13 from customer
14     ,store_sales
15     ,date_dim
16 where c_customer_sk = ss_customer_sk
17     and ss_sold_date_sk = d_date_sk
18 group by c_customer_id
19             ,c_first_name
20             ,c_last_name
21             ,c_preferred_cust_flag
22             ,c_birth_country
23             ,c_login
24             ,c_email_address
25             ,d_year
26 )
27
28 , web_year_total as (
29     select  c_customer_id customer_id
30             ,c_first_name customer_first_name

```

```

31         ,c_last_name customer_last_name
32         ,c_preferred_cust_flag customer_preferred_cust_flag
33         ,c_birth_country customer_birth_country
34         ,c_login customer_login
35         ,c_email_address customer_email_address
36         ,d_year dyear
37         ,sum(ws_ext_list_price-ws_ext_discount_amt) year_total
38 from customer
39     ,web_sales
40     ,date_dim
41 where c_customer_sk = ws_bill_customer_sk
42       and ws_sold_date_sk = d_date_sk
43 group by c_customer_id
44         ,c_first_name
45         ,c_last_name
46         ,c_preferred_cust_flag
47         ,c_birth_country
48         ,c_login
49         ,c_email_address
50         ,d_year
51 )
52
53
54 , t_s_firstyear as (
55     select customer_id
56           ,customer_first_name
57           ,customer_last_name
58           ,customer_email_address
59           ,year_total
60     from store_year_total
61     where dyear = 1999
62     and year_total > 0
63 )
64
65 , t_s_secyear as (
66     select customer_id
67           ,customer_first_name
68           ,customer_last_name
69           ,customer_email_address
70           ,year_total
71     from store_year_total
72     where dyear = 1999+1
73 )
74
75 , t_w_firstyear as (
76     select customer_id
77           ,customer_first_name
78           ,customer_last_name
79           ,customer_email_address
80           ,year_total
81     from web_year_total
82     where dyear = 1999
83     and year_total > 0
84 )
85
86 , t_w_secyear as (
87     select customer_id
88           ,customer_first_name

```

```

89         ,customer_last_name
90         ,customer_email_address
91         ,year_total
92     from web_year_total
93     where dyear = 1999+1
94 )
95
96 select
97     t_s_secyear.customer_id
98     ,t_s_secyear.customer_first_name
99     ,t_s_secyear.customer_last_name
100    ,t_s_secyear.customer_email_address
101 from t_s_firstyear
102 inner join t_s_secyear
103     on t_s_firstyear.customer_id = t_s_secyear.customer_id
104 inner join t_w_firstyear
105     on t_s_firstyear.customer_id = t_w_firstyear.customer_id
106 inner join t_w_secyear
107     on t_w_firstyear.customer_id = t_w_secyear.customer_id
108 where
109     case
110         when t_w_firstyear.year_total > 0 then t_w_secyear.year_total /
111             t_w_firstyear.year_total
112         else 0.0
113     end
114     >
115     case
116         when t_s_firstyear.year_total > 0 then t_s_secyear.year_total /
117             t_s_firstyear.year_total
118         else 0.0
119     end
120 order by t_s_secyear.customer_id
121         ,t_s_secyear.customer_first_name
122         ,t_s_secyear.customer_last_name
123         ,t_s_secyear.customer_email_address
124 limit 100;

```

Listing A.8: Query\_11

```

1  -- Query 74
2
3  with store_year_total as (
4      select c_customer_id customer_id
5             ,c_first_name customer_first_name
6             ,c_last_name customer_last_name
7             ,d_year as year
8             ,stddev_samp(ss_net_paid) year_total
9  from date_dim
10     ,store_sales
11     ,customer
12  where d_date_sk = ss_sold_date_sk
13     and ss_customer_sk = c_customer_sk
14     and d_year in (2001,2001+1)
15  group by c_customer_id
16           ,c_first_name
17           ,c_last_name
18           ,d_year
19 )
20

```

```

21 , web_year_total as (
22   select c_customer_id customer_id
23         ,c_first_name customer_first_name
24         ,c_last_name customer_last_name
25         ,d_year as year
26         ,stddev_samp(ws_net_paid) year_total
27   from date_dim
28        ,web_sales
29        ,customer
30   where d_date_sk = ws_sold_date_sk
31        and c_customer_sk = ws_bill_customer_sk
32        and ws_sold_date_sk = d_date_sk
33        and d_year in (2001,2001+1)
34   group by c_customer_id
35           ,c_first_name
36           ,c_last_name
37           ,d_year
38 )
39
40 , t_s_firstyear as (
41   select customer_id
42         ,customer_first_name
43         ,customer_last_name
44         ,year_total
45   from store_year_total
46   where year = 2001
47         and year_total > 0
48 )
49
50 , t_s_secyear as (
51   select customer_id
52         ,customer_first_name
53         ,customer_last_name
54         ,year_total
55   from store_year_total
56   where year = 2001+1
57 )
58
59 , t_w_firstyear as (
60   select customer_id
61         ,customer_first_name
62         ,customer_last_name
63         ,year_total
64   from web_year_total
65   where year = 2001
66         and year_total > 0
67 )
68
69 , t_w_secyear as (
70   select customer_id
71         ,customer_first_name
72         ,customer_last_name
73         ,year_total
74   from web_year_total
75   where year = 2001+1
76 )
77
78 select

```



```

79 t_s_secyear.customer_id, t_s_secyear.customer_first_name,
   t_s_secyear.customer_last_name
80 from t_s_firstyear
81 inner join t_s_secyear
82   on t_s_firstyear.customer_id = t_s_secyear.customer_id
83 inner join t_w_firstyear
84   on t_s_firstyear.customer_id = t_w_firstyear.customer_id
85 inner join t_w_secyear
86   on t_w_firstyear.customer_id = t_w_secyear.customer_id
87 where
88   (t_w_secyear.year_total / nullif(t_w_firstyear.year_total,0))
89   >
90   (t_s_secyear.year_total / nullif(t_s_firstyear.year_total,0))
91 order by
92   t_s_secyear.customer_last_name
93   ,t_s_secyear.customer_first_name
94   ,t_s_secyear.customer_id
95 limit 100;

```

Listing A.9: Query\_74

```

1  -- Query 81
2
3  with customer_total_return as (
4    select
5      cr_returning_customer_sk as ctr_customer_sk ,
6      ca_state as ctr_state,
7      sum(cr_return_amt_inc_tax) as ctr_total_return
8    from
9      catalog_returns
10     inner join date_dim on cr_returned_date_sk = d_date_sk
11     inner join customer_address on cr_returning_addr_sk =
12       ca_address_sk
13     where
14       d_year = 1998
15     group by
16       cr_returning_customer_sk
17       ,ca_state)
18 , cust_average_return as (
19   select
20     ctr_state as ctr_state
21     ,avg(ctr_total_return) * 1.2 as ctr_avg_return
22   from customer_total_return ctr1
23   group by ctr_state
24 )
25
26 select c_customer_id,
27        c_salutation,
28        c_first_name,
29        c_last_name,
30        ca_street_number,
31        ca_street_name ,
32        ca_street_type,
33        ca_suite_number,
34        ca_city,
35        ca_county,
36        ca_state,
37        ca_zip,

```

```

38     ca_country,
39     ca_gmt_offset ,
40     ca_location_type,
41     ctr_total_return
42 from customer_total_return ctr1
43     inner join cust_average_return ctr2 on ctr1.ctr_state = ctr2.
         ctr_state
44     inner join customer on ctr_customer_sk = c_customer_sk
45     inner join customer_address on c_current_addr_sk = ca_address_sk
46 where
47     ctr1.ctr_total_return > ctr2.ctr_avg_return
48     and ca_state = 'TX'
49 order by c_customer_id,
50     c_salutation,
51     c_first_name,
52     c_last_name,
53     ca_street_number,
54     ca_street_name ,
55     ca_street_type,
56     ca_suite_number,
57     ca_city,
58     ca_county,
59     ca_state,
60     ca_zip,
61     ca_country,
62     ca_gmt_offset ,
63     ca_location_type,
64     ctr_total_return
65 limit 100;

```

Listing A.10: Query\_81